



# Neural Network Optimization for Efficient Inference

Alexander Suslov

# Neural Network Applications

## Self-Driving Car



This image is in the public domain

## Robots



This image is in the public domain

## Image processing




This image is in the public domain

## Machine Translation




This image is in the public domain

## Medical imaging



This image is in the public domain

## 3D scanning



This image is in the public domain

# Where does Inference of Neural Networks Compute?

Standalone



Client-Server

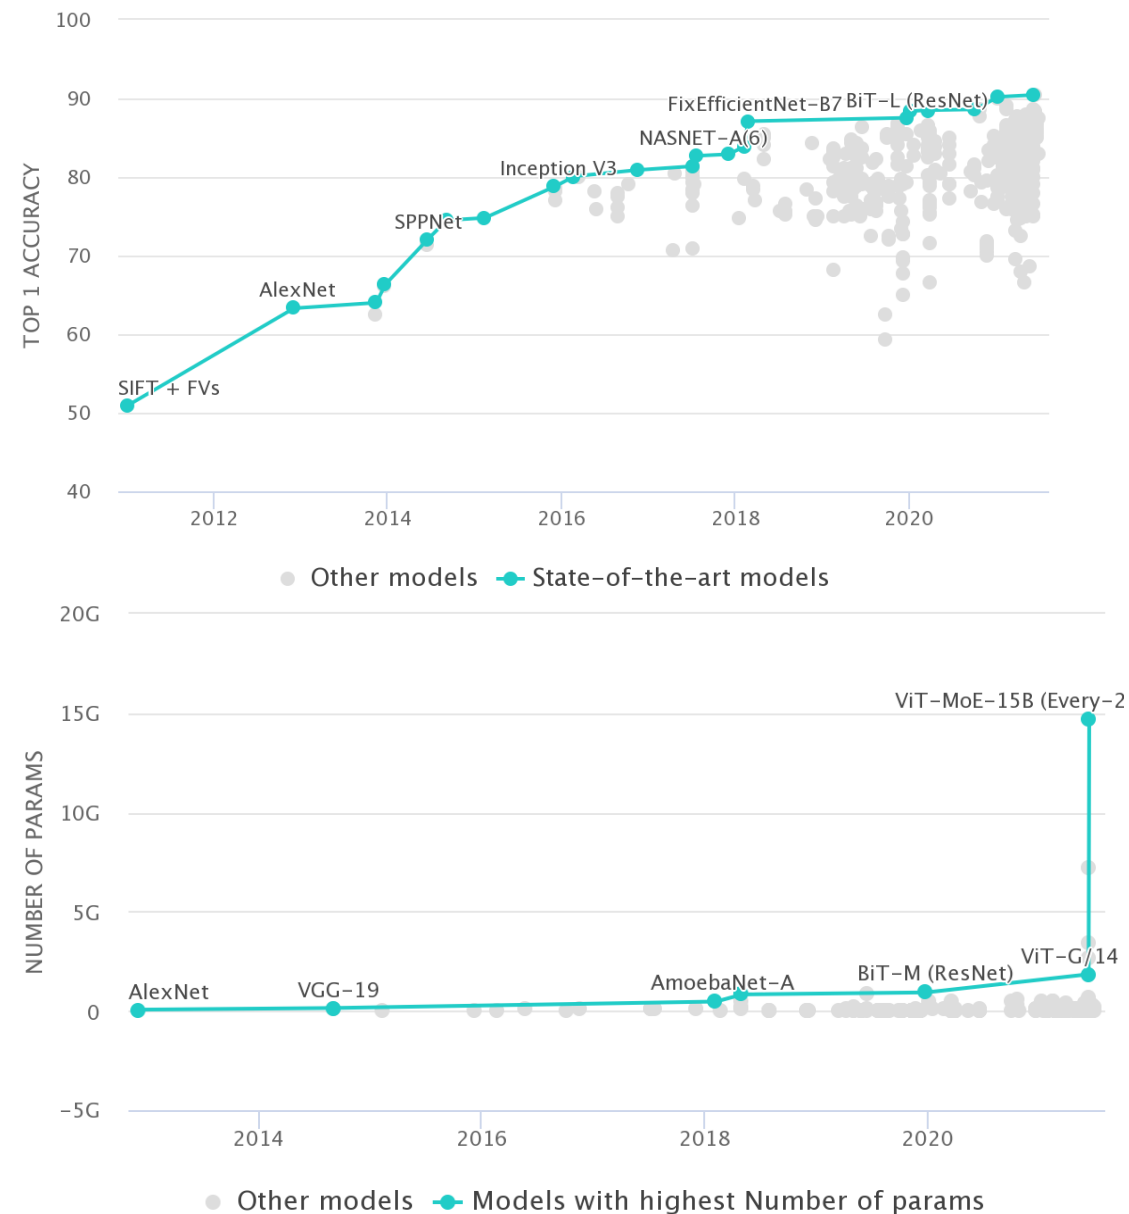


All images are in the public domain

# Models are Getting Larger


- While models are becoming more efficient, high accuracy still implies high complexity

<https://paperswithcode.com/sota/image-classification-on-imagenet>



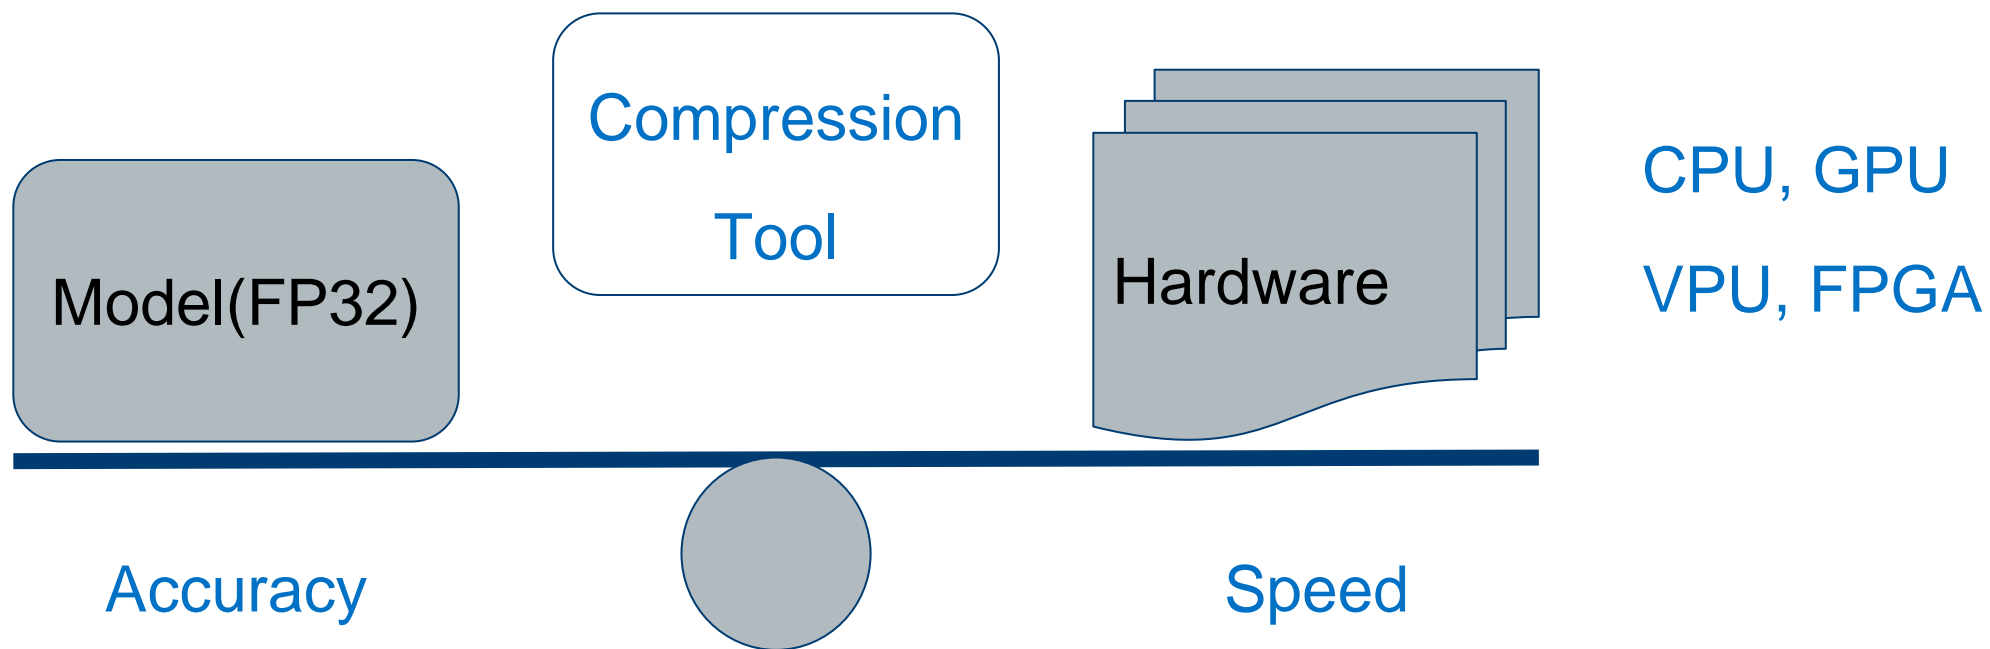
# The First Challenge: Model Size

- Hard to distribute large models through over-the-air update
- The first run is slow due to loading weights.



All images are in the public domain

# The Second Challenge: Speed

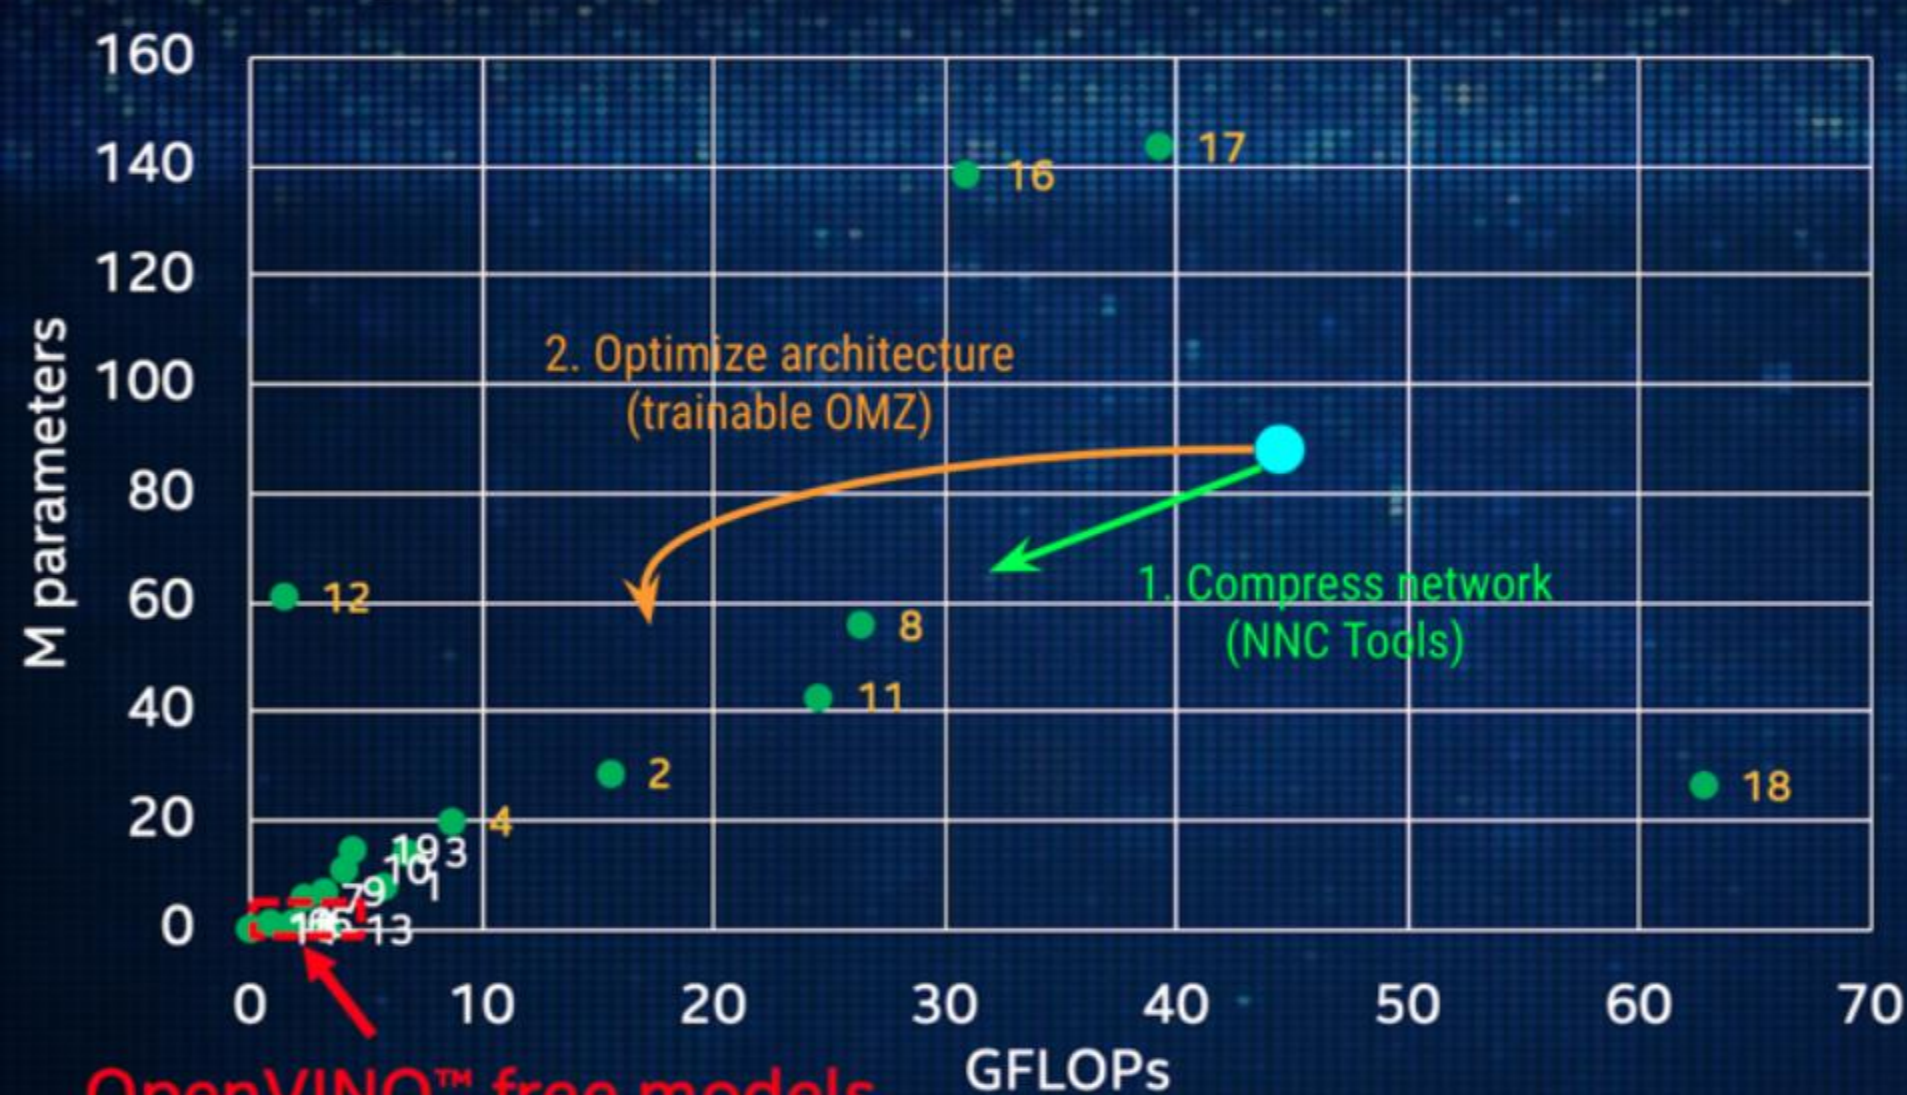


Trade off between accuracy and performance



# Algorithms for Efficient Inference

1. Pruning
2. Weight Clustering
3. Quantization
4. Binary / Ternary Net
5. Distillation
6. Low Rank Approximation



- 1 DenseNet-121
- 2 DenseNet-161
- 3 DenseNet-169
- 4 DenseNet-201
- 5 SqueezeNet1.0
- 6 SqueezeNet1.1
- 7 MobileNet-SSD
- 8 Inception-ResNet-v2
- 9 GoogLeNet-v1
- 10 GoogLeNet-v2
- 11 GoogLeNet-v4
- 12 AlexNet
- 13 MTCNN-p
- 14 MTCNN-r
- 15 MTCNN-o
- 16 VGG16
- 17 VGG19
- 18 SSD300
- 19 MobileNetv2-SSD

OpenVINO™ free models



# Algorithms for Efficient Inference

## 1. Pruning

2. Weight Clustering


3. Quantization

4. Binary / Ternary Net

5. Distillation

6. Low Rank Approximation

# Pruning Neural Networks

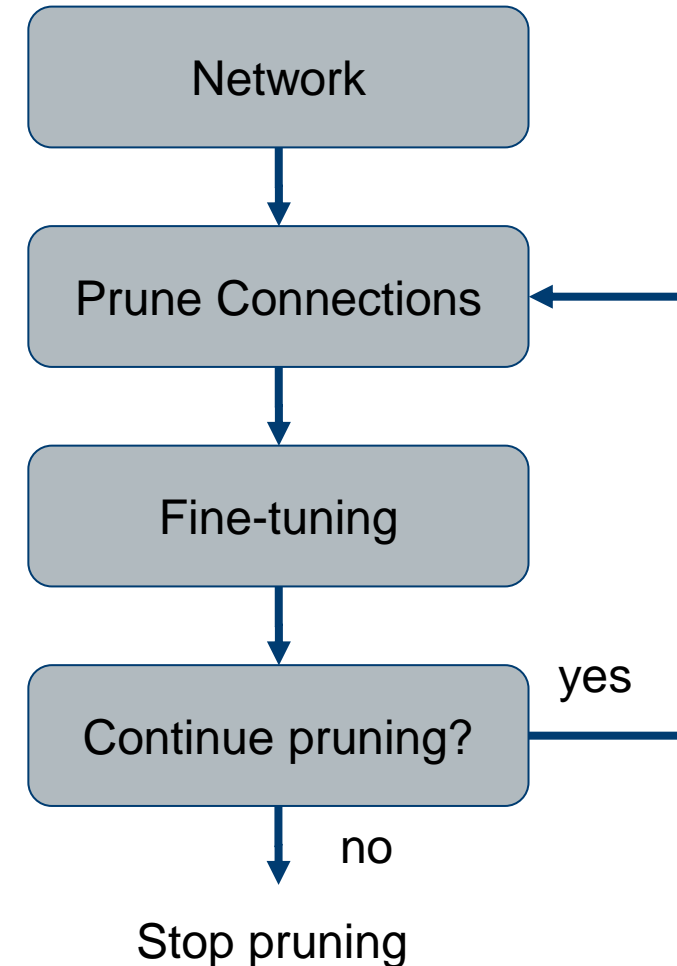


# Pruning Neural Networks

- Neural network pruning benefits
  - Reducing the binary model size
  - Smaller models reduce memory bandwidth bottlenecks
  - Faster kernels (depends on hardware support)



# Pruning Neural Networks

- Criteria for pruning
  - Connections with low weights
  - Neurons or filters with low impact
- External limitations
  - Spatial structure of pruned weights







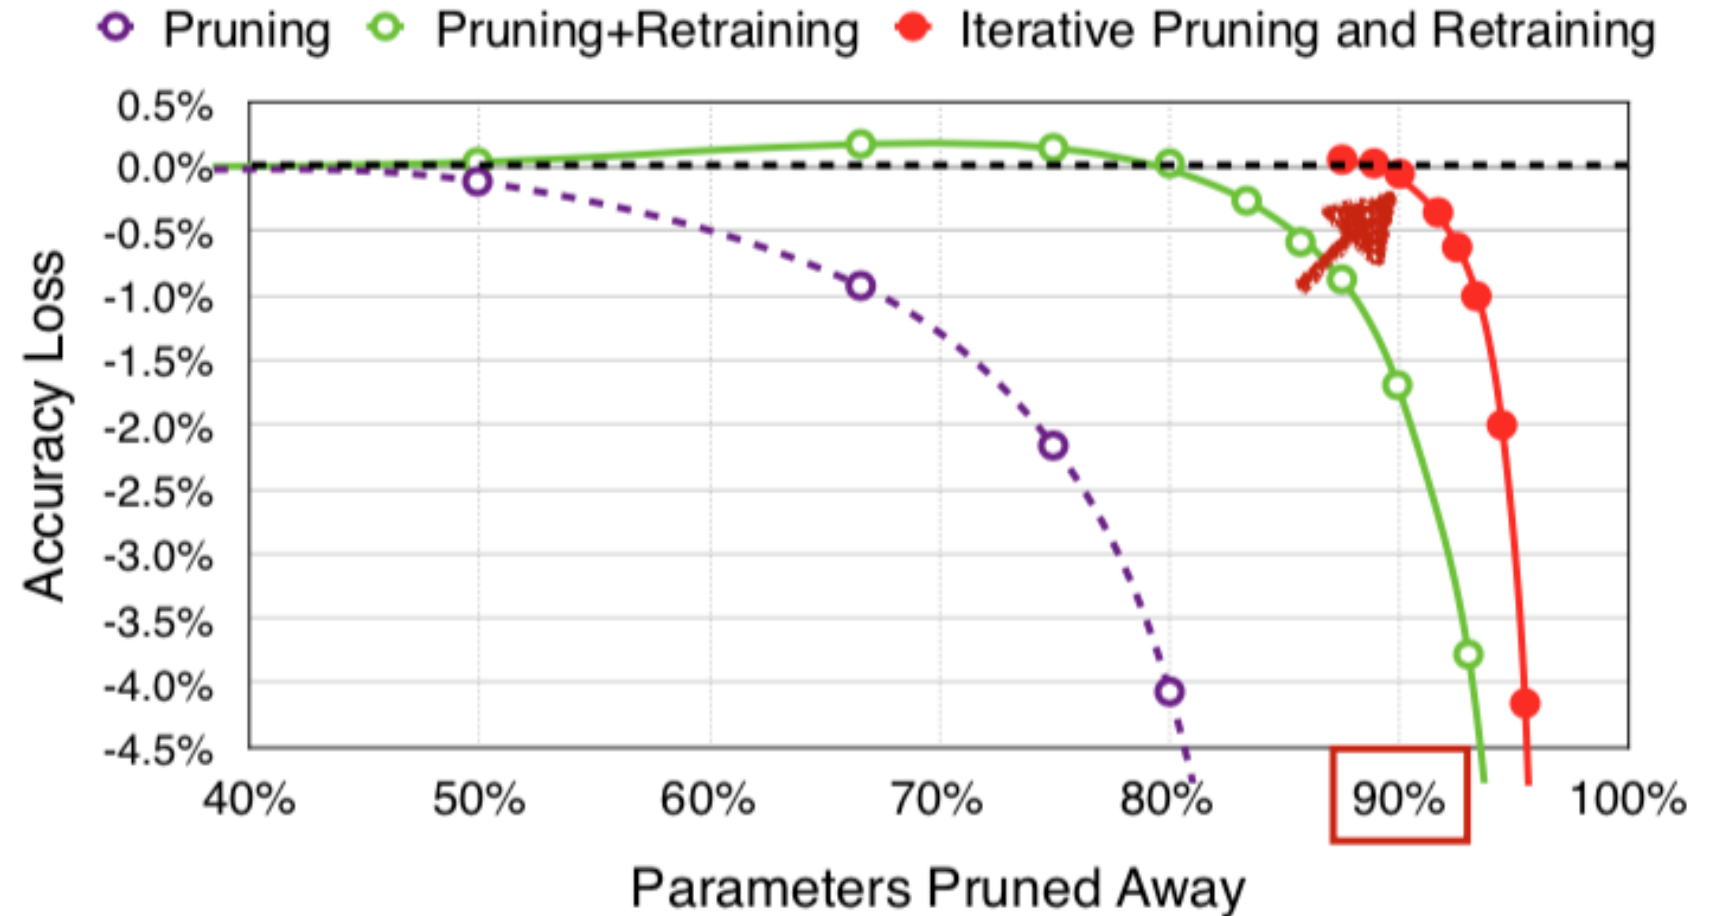

# Pruning Neural Networks



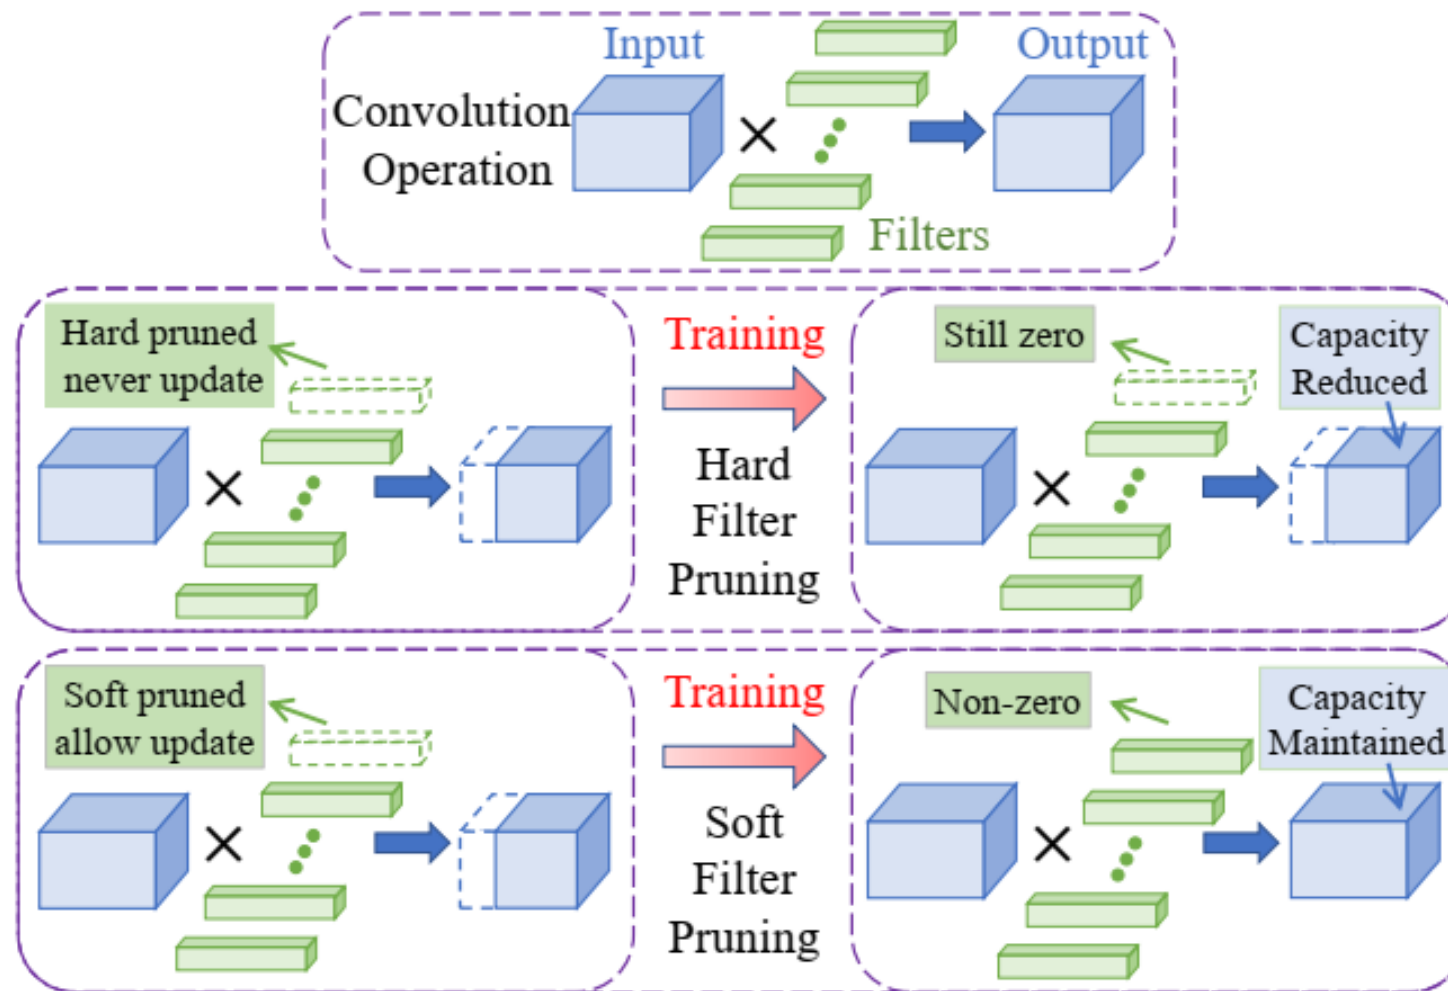
# Retrain to Recover Accuracy



# Iteratively Retrain to Recover Accuracy



# Filter Pruning: Main Approaches





# Filter Pruning: Filter Importance Criteria

- L1, L2

$$\|F\|_p = \sqrt[p]{\sum_{c,k_1,k_2=1}^{C,K,K} |F(c,k_1,k_2)|^p}$$

- “Geometric Median”

$$G(F_i) = \sum_{F_j \in \{F_1, \dots, F_m\}, j \neq i} \|F_i - F_j\|_2$$

- Learned Global Ranking methods

# Filter Pruning: Results

Models	Compression algorithm	Dataset	Top-1 Accuracy FP32 model (%)	Top-1 Accuracy Pruned model (%)
ResNet-18	Filter pruning, 30%, magnitude criterion	ImageNet	69.76	68.69
ResNet-18	Filter pruning, 30%, geometric median criterion	ImageNet	69.76	68.97
ResNet-34	Filter pruning, 30%, magnitude criterion	ImageNet	73.31	72.54
ResNet-34	Filter pruning, 30%, geometric median criterion	ImageNet	73.31	72.60
ResNet-50	Filter pruning, 30%, magnitude criterion	ImageNet	76.13	75.7
ResNet-50	Filter pruning, 30%, geometric median criterion	ImageNet	76.13	75.7

<https://github.com/openvinotoolkit/nncf>

# Sparsification: Main Approaches

- **Magnitude Sparsity**

- After each training epoch the method calculates a threshold based on the current sparsity ratio and uses it to zero weights which are lower than this threshold.

- **Regularization-Based (RB) Sparsity**

- The sparsification algorithm based on probabilistic approach and loss regularization.

Ordinary convolution

$$output = conv(x, w)$$

Sparsifying weights we reparametrize weights as:

$$\bar{w} = w * z$$

Where:

$w$  – weights,  $w \in R$

$z$  – binary mask,  $z \in [0, 1]$

We train these masks using modified loss

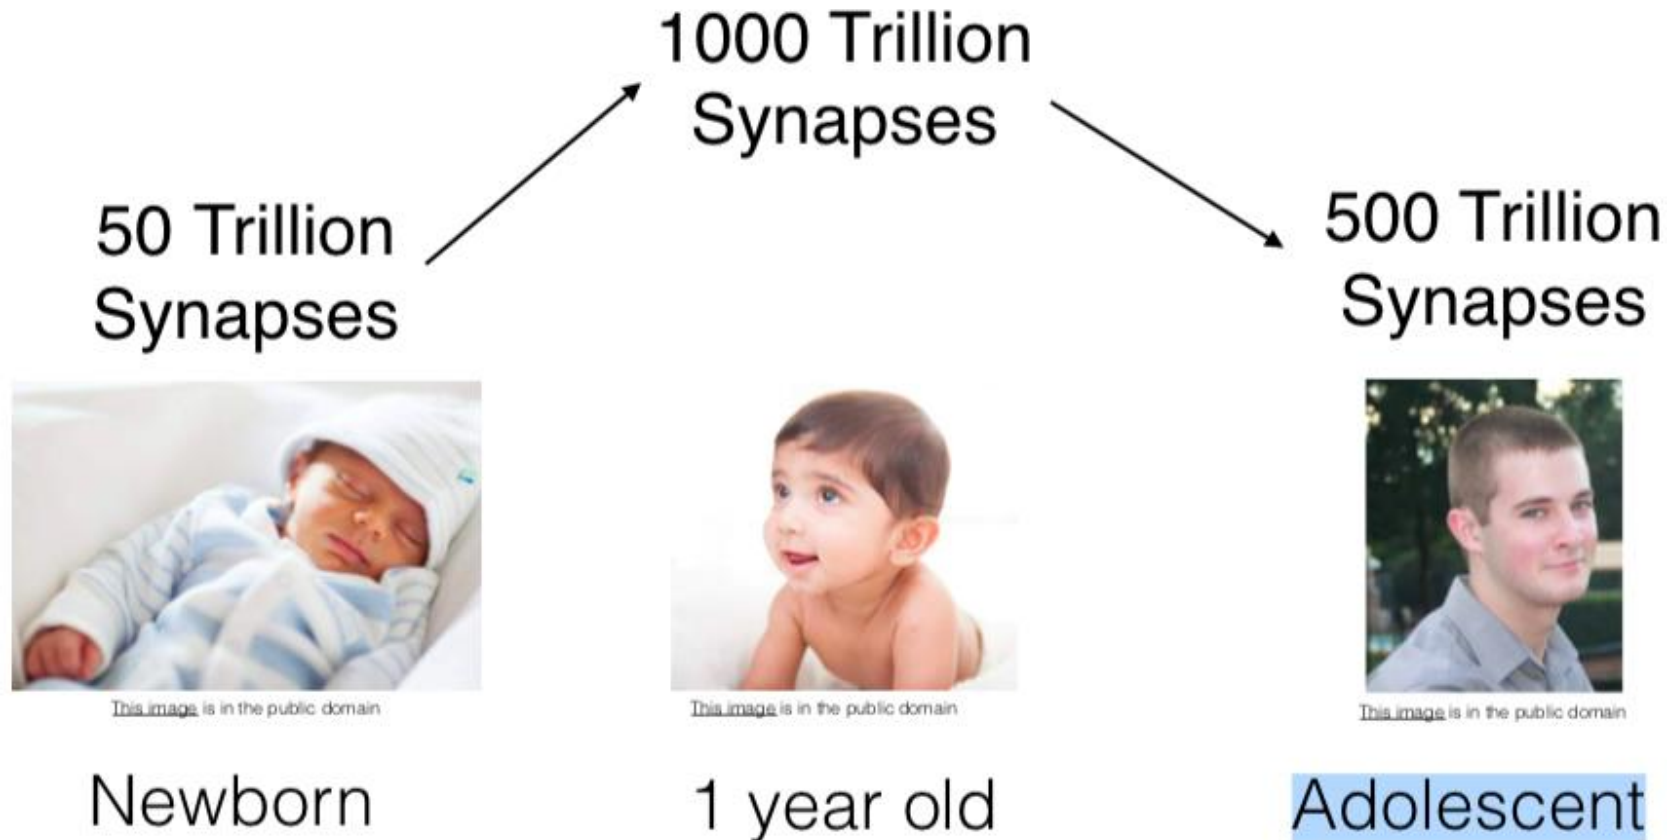
$$Loss = Loss_{task} + \alpha \left( \sum_l^{Layers} ||z|| - target \right)^2$$

# Sparsification: Results

Models	Compression algorithm	Dataset	Top-1 Accuracy FP32 model (%)	Top-1 Accuracy Pruned model (%)
inception-v3	RB-sparsity, 50% sparsity rate	ImageNet	77.46	77.25
inception-v3	Magnitude sparsity, 50% sparsity rate	ImageNet	77.46	77.24
inception-v3	RB-sparsity, 92% sparsity rate	ImageNet	77.46	76.6
mobilenet-v2	RB-sparsity, 50% sparsity rate	ImageNet	71.8	71.2
mobilenet-v2	Magnitude sparsity, 50% sparsity rate	ImageNet	71.8	70.8
mobilenet-v2	RB-sparsity, 78% sparsity rate	ImageNet	71.8	69.98



# Pruning Happens in Human Brain




Christopher A Walsh. Peter Huttenlocher (1931-2013). *Nature*, 502(7470):172–172, 2013.

Slide credits by Song Han

# Algorithms for Efficient Inference

1. Pruning
2. Weight Clustering
3. Quantization
4. Binary / Ternary Net
5. Distillation
6. Low Rank Approximation

# Weight sharing



# Pruning + Weight sharing + Huffman Encoding

Network	Original Size	Compressed Size	Compression Ratio	Original Accuracy	Compressed Accuracy
LeNet-300	1070KB	→ 27KB	<b>40x</b>	98.36%	→ 98.42%
LeNet-5	1720KB	→ 44KB	<b>39x</b>	99.20%	→ 99.26%
AlexNet	240MB	→ 6.9MB	<b>35x</b>	80.27%	→ 80.30%
VGGNet	550MB	→ 11.3MB	<b>49x</b>	88.68%	→ 89.09%
GoogleNet	28MB	→ 2.8MB	<b>10x</b>	88.90%	→ 88.92%
ResNet-18	44.6MB	→ 4.0MB	<b>11x</b>	89.24%	→ 89.28%




# Algorithms for Efficient Inference

1. Pruning
2. Weight Clustering
- 3. Quantization**
4. Binary / Ternary Net
5. Distillation
6. Low Rank Approximation

# Quantization: Overview

This is the process of transforming a neural network such that it can be represented and executed at a lower precision by discretizing the original neural network weights and activations.




# Quantization: Overview


- **Quantization function** is a mapping of values from high to low precision
- **Neural network transformation** is the process of getting  $g'$  from  $g$
- **Quantization algorithm** computes quantization parameters required by new neural network  $g'$  and optimizes  $g'$  via fine-tuning or post-training algorithms.
  - Post-training quantization without dataset ([POT OpenVINO](#))
  - Post-training quantization with dataset ([POT OpenVINO](#))
  - Quantization aware training ([NNCF OpenVINO](#))

# Quantization: Quantization function


- Quantization refers to mapping values from fp32 to a lower precision format with specified parameters:
  - Precision
  - Quantization type
  - Granularity



# Quantization: Quantization types



Asymmetric Mode



Symmetric Mode

# Asymmetric Quantization

- Quantization:

$$x_{int} = \text{round}\left(\frac{x}{\Delta}\right) + z$$

$$x_Q = \text{clamp}(0, N_{levels} - 1, x_{int})$$

- $\Delta$  specifies the step size of the quantizer and floating point zero maps to zero-point.
- $z$  - zero-point.
- $N_{levels} = 256$  for 8-bits of precision

- De-quantization:

$$x_{float} = (x_Q - z)\Delta$$

# Symmetric Quantization

- Quantization, zero-point = 0

- Activations:

$$x_{int} = \text{round}\left(\frac{x}{\Delta}\right)$$

$$x_Q = \text{clamp}(-N_{levels}/2, N_{levels}/2 - 1, x_{int}) \quad \text{if signed}$$

$$x_Q = \text{clamp}(0, N_{levels} - 1, x_{int}) \quad \text{if un-signed}$$

- Weights

$$x_Q = \text{clamp}(-(N_{levels}/2 - 1), N_{levels}/2 - 1, x_{int}) \quad \text{if signed}$$

$$x_Q = \text{clamp}(0, N_{levels} - 2, x_{int}) \quad \text{if un-signed}$$



# Calculations with symmetric vs asymmetric

W and X the weight matrix and input for a fully connected layer

## Symmetric

$$\begin{aligned} W \cdot X &\approx s_1(W_{int}) \cdot s_2(X_{int}) \\ &= s_1 s_2 (W_{int} \cdot X_{int}) \end{aligned}$$

Same calculation

## Asymmetric

$$\begin{aligned} W \cdot X &\approx s_1(W_{int} + o_1) \cdot s_2(X_{int} + o_2) \\ &= s_1 s_2 (W_{int} \cdot X_{int}) + s_1 s_2 o_1 X_{int} + s_1 s_2 o_2 W_{int} + s_1 o_1 s_2 o_2 \end{aligned}$$

Unavoidable  
overhead

Precompute, add to  
layer bias

- Symmetric quantization requires less computation
- Symmetric weights/Asymmetric activations are the best option

# Quantization granularity


- We consider different granularities of quantization:
  - Per-layer quantization
    - Same mapping for all elements in a layer.
  - Per-channel quantization:
    - Choose quantizer parameters independently for each row (fc layers) or for each conv kernel (conv layers)

# How to optimize quantization parameters?

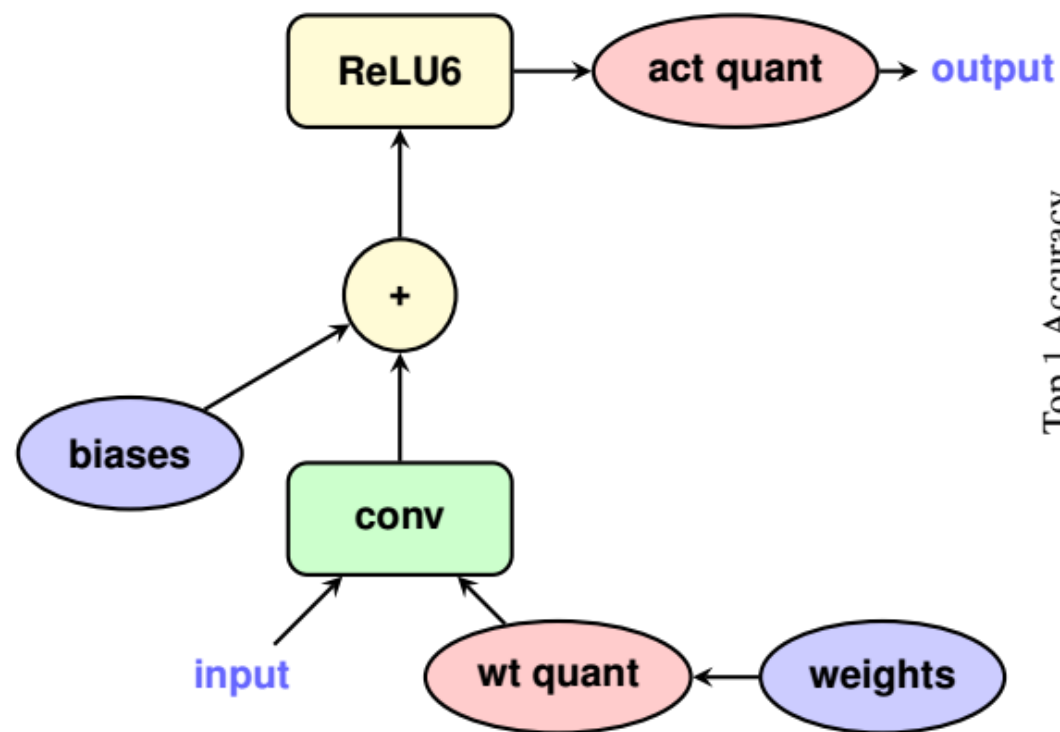
- **Fake-quantization** operation emulates quantization by quantizing and de-quantizing:
  - Values are still in floating point, but with reduced precision

$$\begin{aligned}\text{clamp}(r; a, b) &:= \min(\max(x, a), b) \\ s(a, b, n) &:= \frac{b - a}{n - 1} \\ q(r; a, b, n) &:= \left\lfloor \frac{\text{clamp}(r; a, b) - a}{s(a, b, n)} \right\rfloor s(a, b, n) + a, \\ &\quad (12)\end{aligned}$$

# Quantization: Neural Network Transformation



(a) Integer-arithmetic-only inference



(b) Training with simulated quantization

Top 1 Accuracy

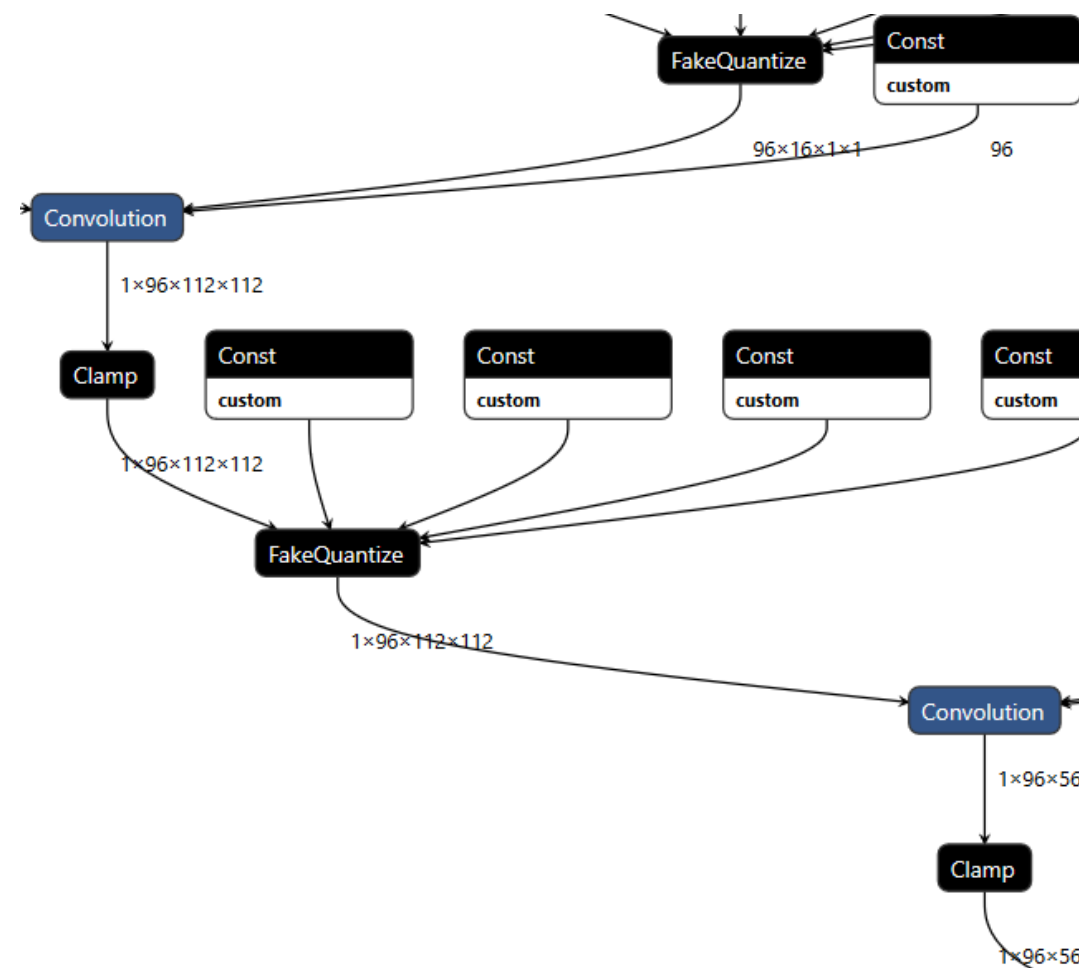
# Quantization: Workflow

Step1: Insert Fake Quantization Layers

Step2: Choose optimal quantization parameters (precision, type, granularity)

Step3: Initialize fake quantization layer parameters

Step4: Optimization

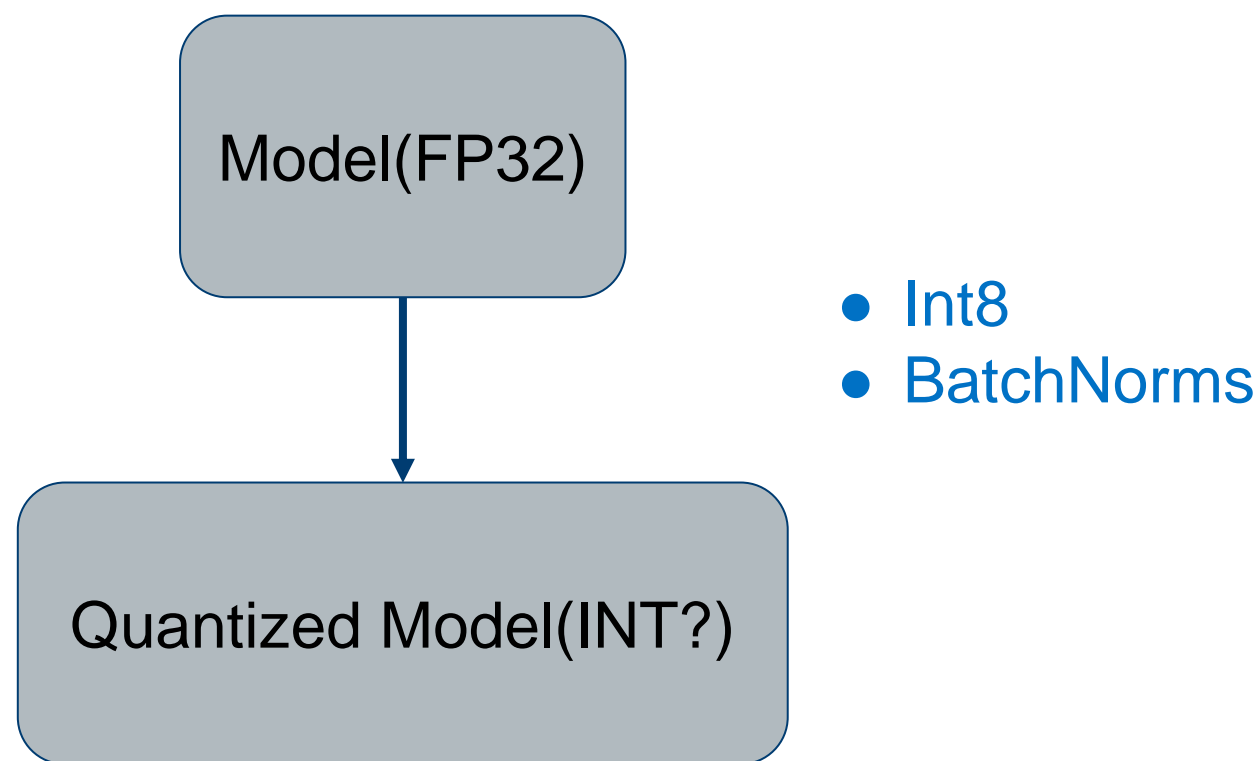


# Quantization algorithms

- Post-Training Quantization without dataset (Data Free Quantization) ([POT OpenVINO](#))
- Post-Training Quantization with dataset ([POT OpenVINO](#))
- Quantization Aware Training ([NNCF OpenVINO](#))

# Quantization algorithms: Usage scenario

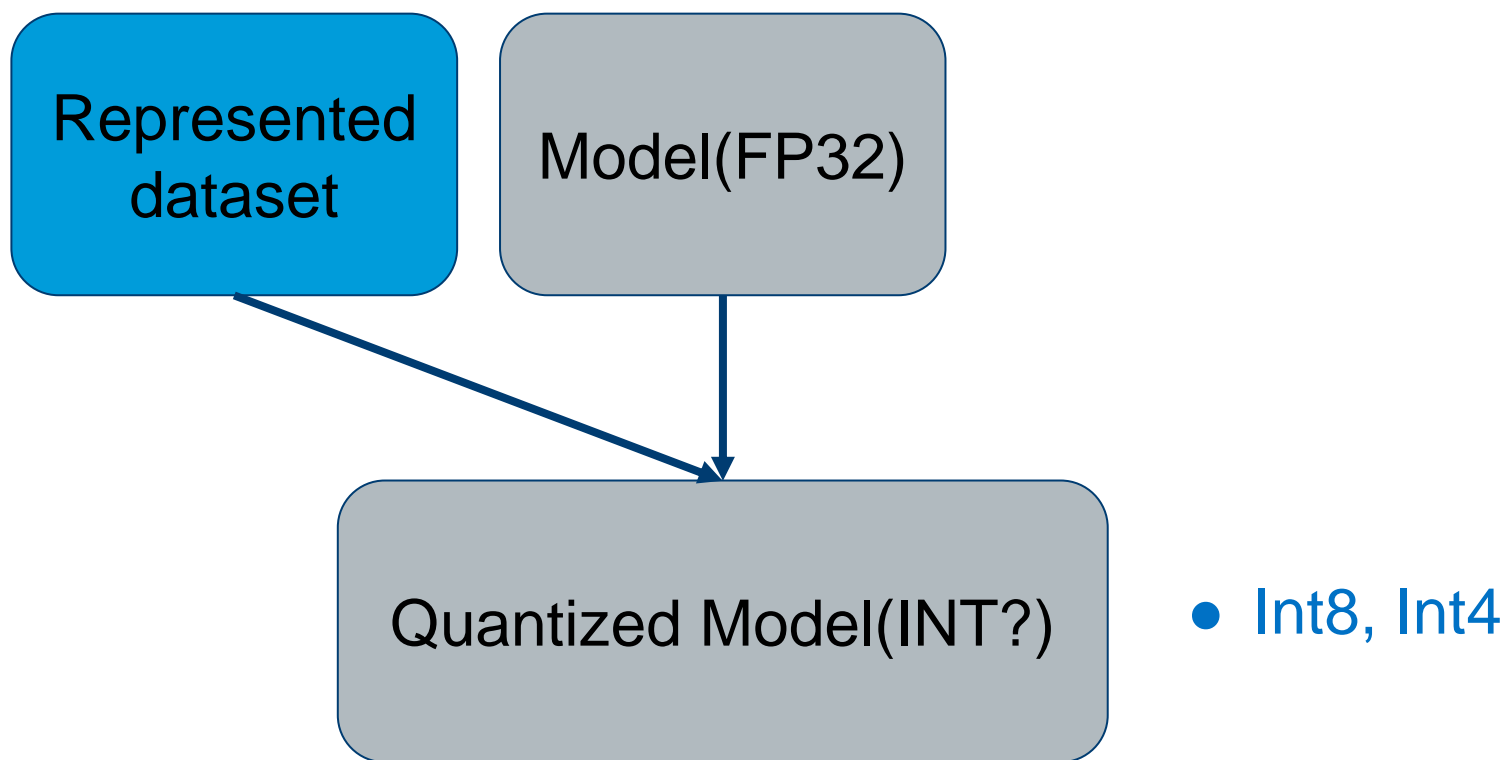
- Level 1: Post-Training Quantization w/o dataset





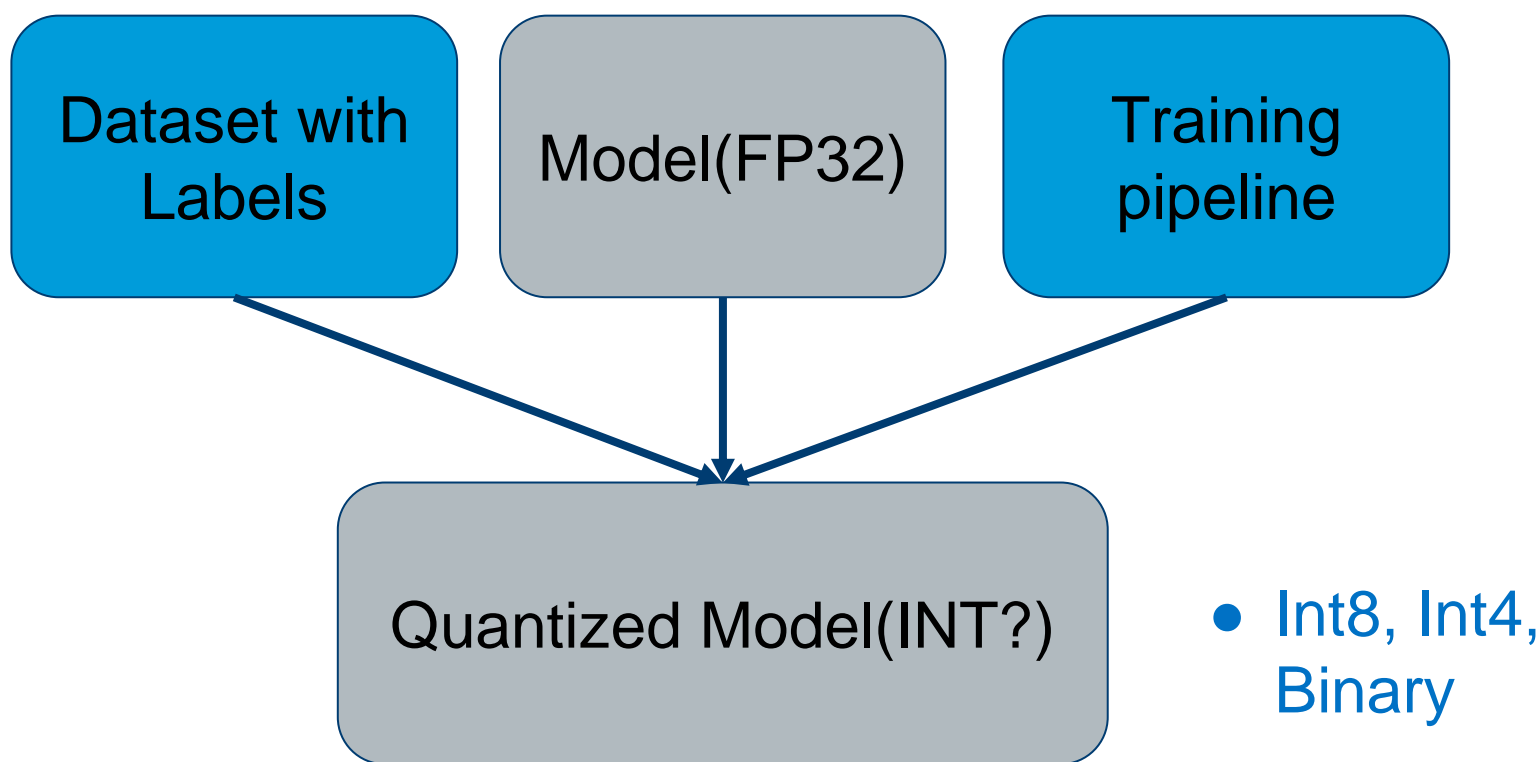
# Quantization algorithms: Usage scenario

- Level 2: Post-Training Quantization w dataset




# Quantization algorithms: Usage scenario

- Level 3: Quantization aware training



# Data-Free Quantization (Level 1)

- The basic idea is to use BatchNorm statistics to estimate the range of activations.
- Algorithm Flow:



Weight tensor  
channel ranges  
alignment

# Quantization with represented dataset (Level 2)

- Weights:
  - $\text{input\_low} = \min(W)$
  - $\text{Input\_high} = \max(W)$
- Activations:
  - Run N examples through the FP32 model and collect for each layer the per-channel pre-activation statistics:
    - moving average of the minimum and maximum values across batches to determine the quantizer parameters for activations.
    - Our observations show that the use of robust statistic boosts accuracy metric (Hodges-Lehmann estimator)

# Bias correction

It is iterative procedure. This procedure is ran on a network with quantized weights only:

1. Run  $N$  examples through the FP32 model and collect for each layer the per-channel pre-activation means  $E[y]$ .
2. For each layer  $L$  in the quantized model:
  - Collect the per-channel pre-activation means  $E[\bar{y}]$  of layer  $L$  for the same  $N$  examples as in step 1.
  - Compute the per-channel biased quantization error  $E[\epsilon] = E[\bar{y}] - E[y]$
  - Subtract  $E[\epsilon]$  from the layer's bias parameter

# Quantization aware training (Level 3)

## Algorithm:

1. Create a training graph of the floating-point model.
2. Insert fake quantization operations in locations where tensors will be downcasted to fewer bits during inference.
3. Train in simulated quantized mode until convergence.
4. Create and optimize the inference graph for running in a low bit inference engine.
5. Run inference using the quantized inference graph.

# Summary: Quantization approaches

	~D	~BP	~AC	MobileNetV2		MobileNetV1		ResNet18		
				FP32	INT8	FP32	INT8	FP32	INT8	INT6
DFQ (ours)	✓	✓	✓	71.7%	<b>71.2%</b>	70.8%	<b>70.5%</b>	69.7%	<b>69.7%</b>	66.3%
Per-layer [18]	✓	✓	✓	71.9%	0.1%	70.9%	0.1%	69.7%	69.2%*	63.8%*
Per-channel [18]	✓	✓	✓	71.9%	69.7%	70.9%	70.3%	69.7%	69.6%*	<b>67.5%*</b>
QT [16] ^	✗	✗	✓	71.9%	70.9%	70.9%	70.0%	-	<b>70.3%</b> <sup>†</sup>	67.3% <sup>†</sup>
SR+DR <sup>†</sup>	✗	✗	✓	-	-	-	71.3%	-	68.2%	59.3%
QMN [31]	✗	✗	✗	-	-	70.8%	68.0%	-	-	-
RQ [21]	✗	✗	✗	-	-	-	<b>70.4%</b>	-	69.9%	<b>68.6%</b>


Table 5. Top1 ImageNet validation results for different models and quantization approaches. The top half compares level 1 approaches (~D: data free, ~BP: backpropagation-free, ~AC: Architecture change free) whereas in the second half we also compare to higher level approaches in literature. Results with \* indicates our own implementation since results are not provided, ^ results provided by [18] and <sup>†</sup> results from table 2 in [21].




# Algorithms for Efficient Inference

1. Pruning
2. Weight Sharing
3. Quantization
- 4. Binary / Ternary Net**
5. Distillation
6. Low Rank Approximation
7. Winograd Transformation

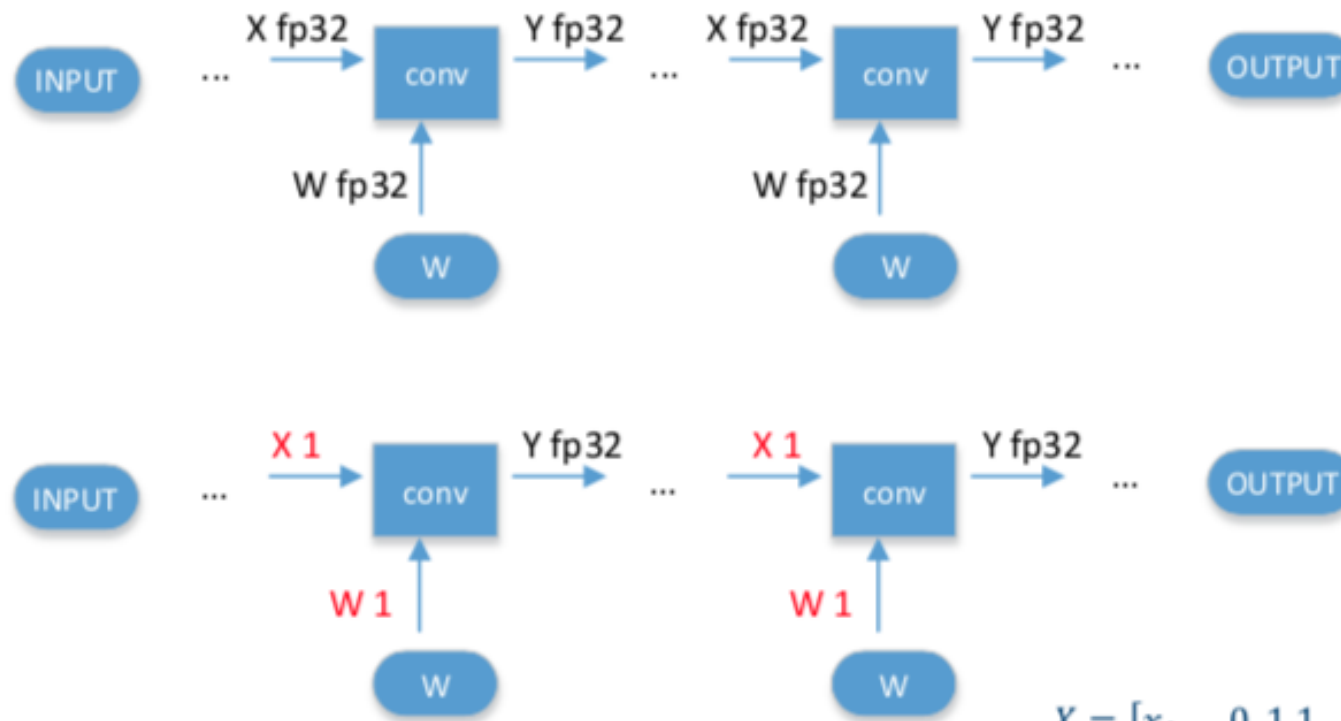
# Binary / Ternary Net



$\Rightarrow$



# Binary Net



1 value inference in conv is dot product  
 $y_1 = w_1 * x_1 + w_2 * x_2 + \dots + w_{32} * x_{32}$   
**32 \* and 32+ operations**

$$y_1 = \{-1, +1\} * \{-1, +1\} + w_2 * x_2 + \dots$$

$$x_i * w_i = \begin{matrix} x \backslash w & -1 & +1 \\ -1 & +1 & -1 \\ +1 & -1 & +1 \end{matrix} \quad \{-1, +1\},$$

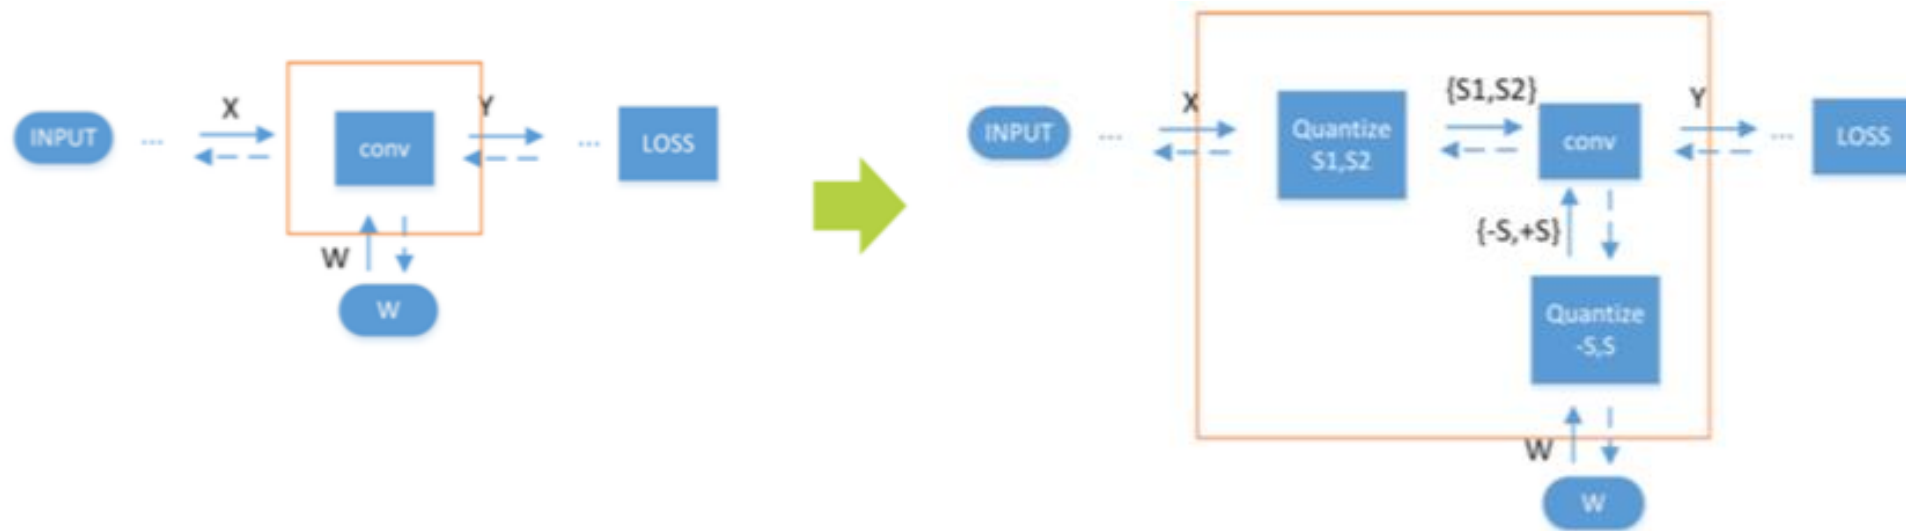
$$a \text{ XNOR } b = \begin{matrix} a \backslash b & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 1 \end{matrix} \quad \{0, 1\},$$

$X = [x_1, \dots, 0, 1, 1, \dots, x_{32}], W = [w_1, \dots, w_{32}] \rightarrow \text{int32 number}$   
 Calc  $y = 2 * \text{popcount}(W \text{ XNOR } X) - 32$   
**We have only 1 XNOR and 1 POPCOUNT operations**

Credited by Konstantin Rodyushkin


# Binary Net

- Convolution layers of the original CNN are replaced with binary convolution alternatives, where inputs activations and weights values are converted into two pretrained values.
- Stage quantization.



Credited by Konstantin Rodyushkin

# Binary Net



Credited by Konstantin Rodyushkin

# Algorithms for Efficient Inference

1. Pruning
2. Weight Sharing
3. Quantization
4. Binary / Ternary Net
- 5. Distillation**
6. Low Rank Approximation

# Distillation

- **L2 – Ba [14]**
  - L2 loss between teacher and student logits
  - No labels required




Figure 6. Teacher-Student model, L2

- [14] J. Ba and R. Caruana, "Do deep nets really need to be deep?" In Advances in neural information processing systems, 2014, pp. 2654–2662.
- [15] G. Hinton, O. Vinyals, and J. Dean, "Distilling the knowledge in a neural network," ArXiv preprint arXiv:1503.02531, 2015.
- [16] A. Romero, N. Ballas, S. E. Kahou, A. Chassang, C. Gatta, and Y. Bengio, "Fitnets: Hints for thin deep nets," ArXiv preprint arXiv:1412.6550, 2015.



# Distillation

- L2 – Ba [14]
  - L2 loss between teacher and student logits
  - No labels required
- **Knowledge Distillation [15]**
  - Soft target: softmax cross entropy with teacher logits
  - Hard target: softmax cross entropy with correct labels




Figure 7. Teacher-Student model, KD

[14] J. Ba and R. Caruana, "Do deep nets really need to be deep?" In Advances in neural information processing systems, 2014, pp. 2654–2662.

[15] G. Hinton, O. Vinyals, and J. Dean, "Distilling the knowledge in a neural network," ArXiv preprint arXiv:1503.02531, 2015.

[16] A. Romero, N. Ballas, S. E. Kahou, A. Chassang, C. Gatta, and Y. Bengio, "Fitnets: Hints for thin deep nets," ArXiv preprint arXiv:1412.6550, 2015.

# Distillation

- L2 – Ba [14]
  - L2 loss between teacher and student logits
  - No labels required
- Knowledge Distillation [15]
  - Soft target: softmax cross entropy with teacher logits
  - Hard target: softmax cross entropy with correct labels
- **FitNets [16]**
  - Knowledge Distillation with hints in the middle points of the network
  - Student is deeper than the teacher




Figure 8. Teacher-Student model, FitNets

[14] J. Ba and R. Caruana, "Do deep nets really need to be deep?" In Advances in neural information processing systems, 2014, pp. 2654–2662.

[15] G. Hinton, O. Vinyals, and J. Dean, "Distilling the knowledge in a neural network," ArXiv preprint arXiv:1503.02531, 2015.

[16] A. Romero, N. Ballas, S. E. Kahou, A. Chassang, C. Gatta, and Y. Bengio, "Fitnets: Hints for thin deep nets," ArXiv preprint arXiv:1412.6550, 2015.

# Algorithms for Efficient Inference

1. Pruning
2. Weight Sharing
3. Quantization
4. Binary / Ternary Net
5. Distillation
6. Low Rank Approximation

# Low Rank Approximation

- Basis filter set => Basis feature maps
- Final feature map = linear combination of basis feature maps
- Rank-1 basis filter => decomposed into a sequence of horizontal and vertical filters
- ~2.4x speedup, no performance drop






Figure 1: (a) The original convolutional layer acting on a single-channel input *i.e.*  $C=1$ . (b) The approximation to that layer using the method of Scheme 1. (c) The approximation to that layer using the method of Scheme 2. Individual filter dimensions are given above the filter layers.



# Summary

Method	Advantages	Disadvantages
Binarization & Quantization	Low latency and memory usage	High loss of accuracy num_bits <= 4
Pruning	Prevents overfitting, the accuracy can increase	Slow
Weight Clustering	Can achieve state of the art results while decreasing the computation cost	Dependent on framework
Distillation	Applicable to all architectures Doesn't change the network	only applicable to classification task

Thank you for your  
attention!